

IAT 461 / 882 — Assignment 2: Modeling Phigma User Behavior

Name: Kayne Hoy SFU ID: 301543404 Date: July 5th, 2026

Before you start: read the full assignment README on Canvas. Run cells top-to-bottom. Do **not** skip cells — later cells depend on earlier ones. Every cell marked `# YOUR CODE HERE` must be completed. Every cell marked `** Your interpretation:**` must be filled in with your own words.

Setup

Run this cell first. It imports all libraries used throughout the notebook.

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.ticker as ticker
import seaborn as sns
import scipy.stats as stats
import statsmodels.api as sm
import statsmodels.formula.api as smf
import warnings
warnings.filterwarnings('ignore')

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import (
    confusion_matrix, classification_report,
    roc_auc_score, roc_curve, ConfusionMatrixDisplay
)
from sklearn.preprocessing import StandardScaler

plt.rcParams.update({
    'figure.dpi': 120,
    'axes.spines.top': False,
    'axes.spines.right': False,
    'axes.titlesize': 13,
```

```
'axes.labelsize': 11,  
)  
ACCENT = '#e8925a'  
  
print("Libraries loaded.")
```

Libraries loaded.

Part A — Linear Regression: Did the AI Assistant Change User Behavior?

Dataset: `phigma_ab_study.csv` **Research question:** Does access to Phigma's AI design assistant increase whiteboard session duration among free-tier users?

This is a **between-subjects experiment**. Users were randomly assigned to treatment (AI access) or control (no AI). We use OLS regression to estimate the effect of the treatment while controlling for other variables.

No train/test split in this part. We are doing statistical inference — explaining the outcome — not predicting future users.

Load the data

```
In [2]: df_ab = pd.read_csv('phigma_ab_study.csv', keep_default_na=False)  
  
print(f"Shape: {df_ab.shape}")  
print(f"\nColumn types:\n{df_ab.dtypes}")  
print(f"\nFirst 3 rows:")  
df_ab.head(3)
```

Shape: (2000, 12)

Column types:

user_id	object
ab_group	object
device_type	object
account_age_days	int64
prior_whiteboard_sessions	int64
plan_region	object
num_collaborators	int64
feature_tour_completed	int64
templates_used	int64
ai_suggestions_used	int64
whiteboard_actions_per_session	int64
avg_whiteboard_session_duration_min	float64
dtype:	object

First 3 rows:

```
Out[2]:
```

	user_id	ab_group	device_type	account_age_days	prior_whiteboard_sessions	plan_region	num_collaborators	feature_tour_compl
0	U76582	control	tablet	478	4	EU	3	
1	U49158	control	desktop	633	9	APAC	1	
2	U82668	control	desktop	159	10	NA	5	

```
In [3]: # Summary statistics
df_ab.describe(include='all').round(2)
```

Out [3]:	user_id	ab_group	device_type	account_age_days	prior_whiteboard_sessions	plan_region	num_collaborators	feature_tour_c
	count	2000	2000	2000	2000.00	2000.00	2000	2000.00
	unique	1978	2	3	NaN	NaN	4	NaN
	top	U79488	control	desktop	NaN	NaN	NA	NaN
	freq	2	1011	1211	NaN	NaN	677	NaN
	mean	NaN	NaN	NaN	356.43	4.51	NaN	2.16
	std	NaN	NaN	NaN	207.82	3.33	NaN	2.37
	min	NaN	NaN	NaN	7.00	0.00	NaN	0.00
	25%	NaN	NaN	NaN	178.00	2.00	NaN	0.00
	50%	NaN	NaN	NaN	352.50	4.00	NaN	1.00
	75%	NaN	NaN	NaN	533.25	6.00	NaN	3.00
	max	NaN	NaN	NaN	729.00	22.00	NaN	15.00

```
In [4]: # Missing values
df_ab.isnull().sum()
```

Out [4]:

	0
user_id	0
ab_group	0
device_type	0
account_age_days	0
prior_whiteboard_sessions	0
plan_region	0
num_collaborators	0
feature_tour_completed	0
templates_used	0
ai_suggestions_used	0
whiteboard_actions_per_session	0
avg_whiteboard_session_duration_min	0

dtype: int64

A1 — Exploratory Data Analysis

A1.1 — Outcome distribution

```
In [5]: # This cell checks the outcome distribution before modeling.
# I compare the raw outcome with the log-transformed outcome because the
# lecture showed that right-skewed variables can become easier to model after a log transform.

# Target outcome variable
outcome = 'avg_whiteboard_session_duration_min'
y_raw = df_ab[outcome]
df_ab['log_duration'] = np.log(df_ab[outcome])
y_log = df_ab['log_duration']

fig, axes = plt.subplots(1, 4, figsize=(20, 4))

# Raw outcome histogram
axes[0].hist(y_raw, bins=30, edgecolor='black', color=ACCENT)
```

```

axes[0].set_title('Raw session duration')
axes[0].set_xlabel('Average session duration (min)')
axes[0].set_ylabel('Number of users')

# Q-Q plot for the raw outcome
# This shows whether the raw outcome follows the normal distribution closely.
stats.probplot(y_raw, dist='norm', plot=axes[1])
axes[1].set_title('Q-Q plot of raw duration')
axes[1].set_xlabel('Theoretical quantiles')
axes[1].set_ylabel('Ordered values')


# Log outcome histogram
axes[2].hist(y_log, bins=30, edgecolor='black', color=ACCENT)
axes[2].set_title('Log session duration')
axes[2].set_xlabel('log(duration)')
axes[2].set_ylabel('Number of users')

# Q-Q plot for the log outcome
# I use this to check whether the log transform improves the shape.
stats.probplot(y_log, dist='norm', plot=axes[3])
axes[3].set_title('Q-Q plot of log(duration)')
axes[3].set_xlabel('Theoretical quantiles')
axes[3].set_ylabel('Ordered values')

plt.tight_layout()
plt.show()

print(f"Raw skewness: {y_raw.skew():.3f}")
print(f"Log skewness: {y_log.skew():.3f}")

```



Raw skewness: 1.377

Log skewness: -0.213

 **Your interpretation (A1.1):**

Original prompt: Is the log-transformed outcome more normally distributed? How does the Q-Q plot support your conclusion? Will you use the raw or log outcome in your model?

My interpretation:

The raw average session duration is not normally distributed because it is clearly right-skewed. Most users have shorter or moderate sessions, while a smaller group has very long sessions, so the right tail is pulled outward. The raw Q-Q plot also supports this because the points move away from the reference line, especially at the upper end, which means the raw outcome does not follow a normal shape very closely. After I apply the log transformation, the histogram becomes more balanced and the skewness changes from about 1.38 to about -0.21, which is much closer to zero. The Q-Q plot for `log_duration` is still not perfect at the two tails, but the middle points follow the reference line more closely than the raw outcome. This connects to the regression lecture because OLS does not require the outcome itself to be perfectly normal, but the model is easier to interpret and the residual assumptions are more reasonable when the outcome is less extremely skewed. The lecture also emphasized that transformations should have a reason, not just be applied automatically. In this case, the log transformation has a clear reason: it reduces the effect of very long sessions and makes differences easier to read as approximate percentage changes. Because of this, I will use `log_duration` in the OLS model.

A1.2 — Examine the confounders

```
In [6]: # This cell focuses on two variables that are measured during the experiment window.
# They are important to inspect, but they are problematic as controls because they happen
# after treatment assignment and are very close to the outcome itself.


fig, axes = plt.subplots(1, 2, figsize=(12, 4))

# AI suggestions are only available to the treatment group, so this variable partly encodes treatment itself.
sns.boxplot(data=df_ab, x='ab_group', y='ai_suggestions_used', ax=axes[0], color=ACCENT)
axes[0].set_title('AI suggestions used by group')
axes[0].set_xlabel('A/B group')
axes[0].set_ylabel('AI suggestions used')

# Actions per session is almost part of the same behavior as session duration.
axes[1].scatter(
    df_ab['whiteboard_actions_per_session'],
    df_ab['log_duration'],
    alpha=0.35,
    s=14,
    color='darkblue'
)
axes[1].set_title('Actions per session vs log(duration)')
axes[1].set_xlabel('Whiteboard actions per session')
axes[1].set_ylabel('log(session duration)')
```

```
plt.tight_layout()
plt.show()

corr_actions = df_ab['whiteboard_actions_per_session'].corr(df_ab['log_duration'])
print(f"Correlation between whiteboard actions/session and log(duration): {corr_actions:.3f}")
print("\nAI suggestions used by group:")
print(df_ab.groupby('ab_group')['ai_suggestions_used'].describe().round(2))
```



Correlation between whiteboard actions/session and log(duration): 0.914

AI suggestions used by group:

	count	mean	std	min	25%	50%	75%	max
ab_group								
control	1011.0	0.00	0.00	0.0	0.0	0.0	0.0	0.0
treatment	989.0	7.33	4.55	0.0	4.0	7.0	10.0	29.0

 **Your interpretation (A1.2):**

Original prompt: What do you notice about `ai_suggestions_used` across the two groups? What does the correlation between `whiteboard_actions_per_session` and the outcome suggest? Why should neither be included in your model?

My interpretation:

`ai_suggestions_used` is zero for the control group because those users did not have access to the AI assistant. This means it is not a clean baseline predictor; it is partly produced by the treatment itself. If I included it in the model, I would be controlling for something that

only exists after treatment assignment, which would make the treatment coefficient harder to interpret.

`whiteboard_actions_per_session` also should not be included because it is very close to the outcome: the correlation with `log_duration` is about `0.91`, and users who stay longer naturally have more time to make more actions. For that reason, I will exclude both variables from the main OLS model because they are post-treatment or outcome-adjacent variables, and including them could control away the effect I am trying to estimate.

A1.3 — Predictor distributions and relationships with outcome

```

In [7]: # This cell checks each numeric predictor with two simple plots:
# 1. its own distribution
# 2. its relationship with log(session duration)

num_predictors = [
    'account_age_days',
    'prior_whiteboard_sessions',
    'num_collaborators',
    'feature_tour_completed',
    'templates_used',
    'ai_suggestions_used',
    'whiteboard_actions_per_session'
]

fig, axes = plt.subplots(len(num_predictors), 2, figsize=(12, 3 * len(num_predictors)))









for i, col in enumerate(num_predictors):
    # Distribution of the predictor
    axes[i, 0].hist(df_ab[col], bins=30, edgecolor='black', color=ACCENT)
    axes[i, 0].set_title(f'Distribution of {col}')
    axes[i, 0].set_xlabel(col)
    axes[i, 0].set_ylabel('Count')









    # Relationship with the log-transformed outcome
    axes[i, 1].scatter(df_ab[col], df_ab['log_duration'], alpha=0.35, s=12, color='darkblue')
    corr = df_ab[col].corr(df_ab['log_duration'])
    axes[i, 1].set_title(f'{col} vs log(duration), r = {corr:.2f}')
    axes[i, 1].set_xlabel(col)
    axes[i, 1].set_ylabel('log(duration)')



plt.tight_layout()
plt.show()

print('Skewness of numeric predictors:')
print(df_ab[num_predictors].skew().sort_values(ascending=False).round(3))

```







Skewness of numeric predictors:

num_collaborators	1.550
ai_suggestions_used	1.439
whiteboard_actions_per_session	1.426
prior_whiteboard_sessions	1.202
templates_used	1.164
account_age_days	0.058
feature_tour_completed	-0.341
dtype:	float64


In [8]: *# This cell compares log session duration across the categorical predictors.
These boxplots help me see whether the groups have different central tendencies.*

```
cat_predictors = ['ab_group', 'device_type', 'plan_region']

fig, axes = plt.subplots(1, len(cat_predictors), figsize=(14, 4))

for i, col in enumerate(cat_predictors):
    sns.boxplot(data=df_ab, x=col, y='log_duration', ax=axes[i], color=ACCENT)
    axes[i].set_title(f'log(duration) by {col}')
    axes[i].set_xlabel(col)
    axes[i].set_ylabel('log(session duration)' if i == 0 else '')
    axes[i].tick_params(axis='x', rotation=25)

plt.tight_layout()
plt.show()
```



Your interpretation (A1.3):

3–5 sentences summarizing your EDA. Which predictors look most promising? Are any right-skewed and worth log-transforming before use as predictors? Which two features are you excluding and why?

My interpretation:

The EDA suggests that the treatment group has longer average whiteboard sessions than the control group, so `ab_group` is the main explanatory variable I want to test in the OLS model. `prior_whiteboard_sessions` and `num_collaborators` also look promising because they are pre-treatment user/account variables and both have positive relationships with `log_duration`. Several count-like predictors are right-skewed, especially `num_collaborators` and `prior_whiteboard_sessions`, so I will use `log1p()` to reduce the influence of extreme high values while still keeping users with zero values. I will exclude `ai_suggestions_used` and `whiteboard_actions_per_session` because they are measured during the experiment and are not clean baseline controls. I will also exclude `templates_used` from the main treatment-effect model because it is measured during the experiment window, so it may be part of the behavior caused by the treatment rather than a fair control variable.

A2 — Feature Engineering

Prepare your feature matrix for OLS. Exclude the two problematic variables identified in A1.

```
In [9]: # Load baseline dataset

# // 1. Encode Categorical Variables
# Convert experimental group ('ab_group') into a binary integer indicator
```

```

# Generate dummy indicators (one-hot encoding) for device types, omitting the first category
# Generate dummy indicators (one-hot encoding) for regions, omitting the first category

# // 2. Perform Optional Transformations
# Log-transform right-skewed numerical predictors using log(x + 1) to avoid zero values

# // 3. Assemble Modeling Dataset
# Merge numerical characteristics, binary groups, and dummy indicators into matrix X
# Set target vector y as the log-transformed outcome variable

# // 4. Quality Control
# Verify that there are no missing (null) values in feature matrix X
# Inspect matrix dimensions and print a preview of the processed data

```

```

In [10]: # I build the OLS feature matrix using variables that are available before or at assignment.
# I avoid post-treatment variables so the treatment coefficient is easier to interpret.

# Treatment dummy: 0 = control, 1 = treatment
ab_dummy = (df_ab['ab_group'] == 'treatment').astype(int).rename('ab_group')

# Binary onboarding feature
feature_tour = df_ab['feature_tour_completed'].astype(int)

# Log-transform skewed count variables using log1p so zero values are handled safely.
log_account_age = np.log1p(df_ab['account_age_days']).rename('log_account_age_days')
log_prior_sessions = np.log1p(df_ab['prior_whiteboard_sessions']).rename('log_prior_whiteboard_sessions')
log_collaborators = np.log1p(df_ab['num_collaborators']).rename('log_num_collaborators')

# One-hot encode categorical variables and drop one level to avoid the dummy variable trap.
device_dummies = pd.get_dummies(df_ab['device_type'], prefix='device_type', drop_first=True, dtype=int)
region_dummies = pd.get_dummies(df_ab['plan_region'], prefix='plan_region', drop_first=True, dtype=int)

# Final feature matrix for OLS
X = pd.concat([
    ab_dummy,
    feature_tour,
    log_account_age,
    log_prior_sessions,
    log_collaborators,
    device_dummies,
    region_dummies
], axis=1).astype(float)

# Target is the log-transformed outcome from A1.
y = df_ab['log_duration']

```

```
# Quality checks
print(f"Feature matrix shape: {X.shape}")
print(f"Missing values in X: {X.isna().sum().sum()}")
print(f"Infinite values in X: {np.isinf(X.to_numpy()).sum()}")
print("\nFinal feature columns:")
print(list(X.columns))

X.head(3)
```

Feature matrix shape: (2000, 10)

Missing values in X: 0

Infinite values in X: 0

Final feature columns:

['ab_group', 'feature_tour_completed', 'log_account_age_days', 'log_prior_whiteboard_sessions', 'log_num_collaborators', 'device_type_mobile', 'device_type_tablet', 'plan_region_EU', 'plan_region_LATAM', 'plan_region_NA']

```
Out[10]:
```

	ab_group	feature_tour_completed	log_account_age_days	log_prior_whiteboard_sessions	log_num_collaborators	device_type_mobi
0	0.0	1.0	6.171701	1.609438	1.386294	C
1	0.0	0.0	6.452049	2.302585	0.693147	C
2	0.0	1.0	5.075174	2.397895	1.791759	C

Your interpretation (A2):

List your final features. For each one, write one sentence justifying its inclusion and any transformation applied.

My interpretation:

My final OLS feature set is `ab_group`, `feature_tour_completed`, `log_account_age_days`, `log_prior_whiteboard_sessions`, `log_num_collaborators`, the device dummy variables, and the region dummy variables. I included `ab_group` because it is the experimental treatment variable, so it directly answers the question of whether AI assistant access changed session duration. I included `feature_tour_completed` because onboarding may affect how comfortable users are with the whiteboard tool and may explain some variation in session length. I log-transformed `account_age_days`, `prior_whiteboard_sessions`, and `num_collaborators` because these variables are count-like or skewed. This follows the feature engineering lecture more specifically: when a variable has a long right tail, a few very large values can dominate the scale of the variable and make the model overly influenced by extreme users. Using `log1p()` compresses those high values while still keeping zero values valid, which is important for variables such as prior sessions or collaborators. I included device and region as one-hot encoded control variables because usage patterns may differ by device and geographic region. This also follows the point that unordered categories should not be encoded as simple numbers, because that would

create a fake ranking. I dropped one category from each dummy set to avoid the dummy variable trap, meaning the model has a clear reference group instead of duplicate information. I also checked the final feature matrix and confirmed that it has no missing or infinite values before fitting the model.

A3 — OLS Regression

```
In [11]: # Add a constant for the intercept, then fit the OLS regression model.  
# This follows the idea of using regression as an explanatory model.  
  
X_const = sm.add_constant(X)  
model_ols = sm.OLS(y, X_const).fit()  
print(model_ols.summary())
```

OLS Regression Results

```

=====
Dep. Variable:          log_duration    R-squared:                0.360
Model:                  OLS             Adj. R-squared:           0.357
Method:                 Least Squares   F-statistic:             111.9
Date:                   Mon, 06 Jul 2026 Prob (F-statistic):       1.31e-184
Time:                   06:38:39        Log-Likelihood:          -1025.2
No. Observations:      2000            AIC:                    2072.
Df Residuals:          1989            BIC:                    2134.
Df Model:               10
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	3.3431	0.066	50.975	0.000	3.214	3.472
ab_group	0.2164	0.018	11.924	0.000	0.181	0.252
feature_tour_completed	0.0495	0.018	2.691	0.007	0.013	0.086
log_account_age_days	-0.0171	0.010	-1.703	0.089	-0.037	0.003
log_prior_whiteboard_sessions	0.1420	0.014	10.296	0.000	0.115	0.169
log_num_collaborators	0.0866	0.013	6.856	0.000	0.062	0.111
device_type_mobile	-0.7274	0.027	-27.409	0.000	-0.779	-0.675
device_type_tablet	-0.2585	0.022	-11.979	0.000	-0.301	-0.216
plan_region_EU	0.0752	0.025	3.063	0.002	0.027	0.123
plan_region_LATAM	-0.0105	0.033	-0.323	0.747	-0.075	0.054
plan_region_NA	0.0804	0.024	3.374	0.001	0.034	0.127

```

=====
Omnibus:                0.981    Durbin-Watson:                2.049
Prob(Omnibus):           0.612    Jarque-Bera (JB):           1.035
Skew:                    0.028    Prob(JB):                   0.596
Kurtosis:                2.904    Cond. No.                   45.0
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

A3.1 — Interpret the model

```

In [12]: # This cell extracts the main model statistics and translates log coefficients
# into more understandable percentage changes.

print(f"R-squared:          {model_ols.rsquared:.4f}")
print(f"Adjusted R-squared: {model_ols.rsquared_adj:.4f}")
print()

ab_coef = model_ols.params['ab_group']

```

```

ab_pval = model_ols.pvalues['ab_group']
ab_pct_change = np.exp(ab_coef) - 1
control_typical_duration = np.exp(df_ab.loc[df_ab['ab_group'] == 'control', 'log_duration'].mean())
ab_min_effect = control_typical_duration * ab_pct_change

print('Treatment effect:')
print(f"ab_group coefficient: {ab_coef:.4f}")
print(f"p-value: {ab_pval:.4g}")
print(f"Approx. % change: {ab_pct_change:.1%}")
print(f"Approx. minute change for a typical control user: {ab_min_effect:.2f} minutes")
print()

# Table of coefficients, p-values, and approximate percentage changes.
coef_table = pd.DataFrame({
    'coef_log_scale': model_ols.params,
    'p_value': model_ols.pvalues,
    'approx_percent_change': np.exp(model_ols.params) - 1
}).sort_values('p_value')

print('Coefficient table:')
display(coef_table.round(4))

print('Predictors significant at alpha = 0.05:')
sig = coef_table[(coef_table['p_value'] < 0.05) & (coef_table.index != 'const')]
display(sig.round(4))

```

R-squared: 0.3600

Adjusted R-squared: 0.3567

Treatment effect:

ab_group coefficient: 0.2164

p-value: 1.043e-31

Approx. % change: 24.2%

Approx. minute change for a typical control user: 7.55 minutes

Coefficient table:

	coef_log_scale	p_value	approx_percent_change
const	3.3431	0.0000	27.3064
device_type_mobile	-0.7274	0.0000	-0.5168
device_type_tablet	-0.2585	0.0000	-0.2278
ab_group	0.2164	0.0000	0.2415
log_prior_whiteboard_sessions	0.1420	0.0000	0.1526
log_num_collaborators	0.0866	0.0000	0.0905
plan_region_NA	0.0804	0.0008	0.0837
plan_region_EU	0.0752	0.0022	0.0781
feature_tour_completed	0.0495	0.0072	0.0508
log_account_age_days	-0.0171	0.0887	-0.0170
plan_region_LATAM	-0.0105	0.7471	-0.0105

Predictors significant at $\alpha = 0.05$:

	coef_log_scale	p_value	approx_percent_change
device_type_mobile	-0.7274	0.0000	-0.5168
device_type_tablet	-0.2585	0.0000	-0.2278
ab_group	0.2164	0.0000	0.2415
log_prior_whiteboard_sessions	0.1420	0.0000	0.1526
log_num_collaborators	0.0866	0.0000	0.0905
plan_region_NA	0.0804	0.0008	0.0837
plan_region_EU	0.0752	0.0022	0.0781
feature_tour_completed	0.0495	0.0072	0.0508

Results paragraph (A3):

Write a results paragraph as if reporting in a research paper. Include:

- The treatment effect estimate and its significance

- R^2 and what it tells you about model fit
- At least two other notable predictors and their direction
- Whether the AI assistant's effect is practically meaningful (not just statistically significant)

My results paragraph:

I fitted an OLS regression model predicting `log_duration` from treatment assignment and baseline user characteristics. The model had an R^2 of about `0.360`, meaning it explained about 36% of the variation in log average session duration. This is a meaningful amount of explained variation for user-behavior data, but it also shows that many individual differences are still not captured by the model. The treatment coefficient for `ab_group` was positive and statistically significant (`coef = 0.216` , `p < 0.001`), which means users with AI assistant access had about `24.2%` longer sessions than control users, holding the other variables constant. For a typical control user, this is about a `7.6` minute increase, so I would treat the AI assistant effect as practically meaningful and not only statistically significant. Other notable predictors were device type and prior use: mobile users and tablet users had shorter sessions than desktop users, while users with more prior whiteboard sessions and more collaborators tended to have longer sessions. This result makes logical sense because desktop use may support longer design work, and users with more prior activity or collaborators may already have stronger reasons to spend time in the product.

A3.2 — Regression diagnostics

In [13]: *# Regression diagnostics from the lecture: check residual patterns and residual normality.*

```
fitted = model_ols.fittedvalues
residuals = model_ols.resid


fig, axes = plt.subplots(1, 2, figsize=(12, 4))

# Residuals vs fitted checks whether residual spread changes across predicted values.
axes[0].scatter(fitted, residuals, alpha=0.35, s=14, color='darkblue')
axes[0].axhline(0, color='black', linestyle='--', linewidth=1)
axes[0].set_title('Residuals vs fitted values')
axes[0].set_xlabel('Fitted log(duration)')
axes[0].set_ylabel('Residuals')

# Q-Q plot checks whether errors are approximately normal.
stats.probplot(residuals, dist='norm', plot=axes[1])
axes[1].set_title('Q-Q plot of residuals')

plt.tight_layout()
plt.show()
```

```
print(f"Residual mean: {residuals.mean():.4f}")
print(f"Residual skewness: {residuals.skew():.4f}")
```



Residual mean: 0.0000
Residual skewness: 0.0279

Your interpretation (A3.2):

Do the residuals look randomly scattered around zero? Does the Q-Q plot suggest normally distributed errors? Note any patterns that concern you.

My interpretation:

The residuals are mostly scattered around zero, so I do not see a strong curved pattern that would make the linear model clearly inappropriate. The spread is not perfectly equal across all fitted values, which means the homoscedasticity assumption is not perfect, but the pattern is still reasonable for this assignment after using the log-transformed outcome. The Q-Q plot shows that the middle of the residual distribution follows the normal reference line fairly well. The main concern is at the tails, where the points move away from the line, meaning a few users still have residuals that are larger or smaller than the model expects under a perfect normal distribution. Overall, I would not claim the assumptions are perfectly satisfied, but the diagnostics support using the OLS model as a useful explanatory model.

A4 — Exploratory Interaction Analysis

The average treatment effect hides heterogeneity — the AI assistant may help some users more than others.

 State your hypothesis here (before writing any code):

Which interaction are you testing? Which two variables? What is your theoretical reason for expecting this interaction? Which group do you predict will benefit more from the AI assistant?

My hypothesis before fitting the model:

I will test the interaction between `ab_group` and `log_prior_whiteboard_sessions`. My hypothesis is that the AI assistant may help newer or less experienced users more because they have less prior knowledge of how to use the whiteboard tool efficiently. In other words, I expect the treatment effect to be stronger for users with fewer prior whiteboard sessions and weaker for users who were already active before the experiment. This is theoretically motivated because the AI assistant may act like guidance or support, and guidance should matter more when users do not already know the tool well. This hypothesis also connects to the idea that an average treatment effect can hide different effects across subgroups.

```
In [14]: # Create an interaction between treatment assignment and prior whiteboard experience.

X_interact = X.copy()
X_interact['ab_x_log_prior_whiteboard_sessions'] = (
    X_interact['ab_group'] * X_interact['log_prior_whiteboard_sessions']
)

X_interact_const = sm.add_constant(X_interact)
model_interact = sm.OLS(y, X_interact_const).fit()
print(model_interact.summary())

interaction_terms = pd.DataFrame({
    'coef': model_interact.params,
    'p_value': model_interact.pvalues,
    'approx_percent_change': np.exp(model_interact.params) - 1
}).loc[['ab_group', 'log_prior_whiteboard_sessions', 'ab_x_log_prior_whiteboard_sessions']]

print('\nKey interaction terms:')
display(interaction_terms.round(4))
```

OLS Regression Results

```

=====
Dep. Variable:          log_duration    R-squared:                0.361
Model:                  OLS             Adj. R-squared:           0.357
Method:                 Least Squares    F-statistic:             102.0
Date:                  Mon, 06 Jul 2026  Prob (F-statistic):       3.64e-184
Time:                  06:38:39          Log-Likelihood:          -1023.8
No. Observations:      2000             AIC:                    2072.
Df Residuals:          1988             BIC:                    2139.
Df Model:              11
Covariance Type:       nonrobust
=====

```

```

=====
               coef      std err          t      P>|t|      [0.025      0.975]
-----
const                3.3734      0.068     49.564      0.000        3.240        3.507
ab_group              0.1471      0.046      3.231      0.001        0.058        0.236
feature_tour_completed 0.0495      0.018      2.687      0.007        0.013        0.086
log_account_age_days  -0.0164      0.010     -1.636      0.102       -0.036        0.003
log_prior_whiteboard_sessions 0.1194      0.019      6.165      0.000        0.081        0.157
log_num_collaborators  0.0869      0.013      6.881      0.000        0.062        0.112
device_type_mobile    -0.7265      0.027    -27.382      0.000       -0.779       -0.674
device_type_tablet    -0.2587      0.022    -11.992      0.000       -0.301       -0.216
plan_region_EU         0.0756      0.025      3.081      0.002        0.027        0.124
plan_region_LATAM     -0.0111      0.033     -0.340      0.734       -0.075        0.053
plan_region_NA         0.0792      0.024      3.326      0.001        0.032        0.126
ab_x_log_prior_whiteboard_sessions 0.0457      0.028      1.658      0.098       -0.008        0.100
=====

```

```

=====
Omnibus:                1.324    Durbin-Watson:                2.050
Prob(Omnibus):          0.516    Jarque-Bera (JB):            1.348
Skew:                   0.028    Prob(JB):                    0.510
Kurtosis:               2.886    Cond. No.                     49.5
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Key interaction terms:

	coef	p_value	approx_percent_change
ab_group	0.1471	0.0013	0.1585
log_prior_whiteboard_sessions	0.1194	0.0000	0.1268
ab_x_log_prior_whiteboard_sessions	0.0457	0.0975	0.0468


```

In [15]: # Visualize the interaction by grouping users into low, medium, and high prior-use groups.
# This makes the continuous interaction easier to read.

plot_df = df_ab.copy()
plot_df['prior_session_group'] = pd.qcut(
    plot_df['prior_whiteboard_sessions'],
    q=3,
    labels=['Low prior use', 'Medium prior use', 'High prior use']
)


mean_duration = (
    plot_df
    .groupby(['prior_session_group', 'ab_group'])['avg_whiteboard_session_duration_min']
    .mean()
    .reset_index()
)

fig, ax = plt.subplots(figsize=(8, 5))
sns.barplot(
    data=mean_duration,
    x='prior_session_group',
    y='avg_whiteboard_session_duration_min',
    hue='ab_group',
    ax=ax
)
ax.set_title('Mean session duration by treatment and prior use')
ax.set_xlabel('Prior whiteboard session group')
ax.set_ylabel('Average session duration (min)')
ax.legend(title='A/B group')
plt.tight_layout()
plt.show()

mean_duration

```

Mean session duration by treatment and prior use



Out[15]:

	prior_session_group	ab_group	avg_whiteboard_session_duration_min
--	---------------------	----------	-------------------------------------

0	Low prior use	control	32.681935
1	Low prior use	treatment	39.315618
2	Medium prior use	control	36.179009
3	Medium prior use	treatment	44.192039
4	High prior use	control	37.995235
5	High prior use	treatment	50.423934

Your interpretation (A4):

Was your hypothesis supported? What does the interaction coefficient tell you about which users benefit most from the AI assistant? What does this imply for how Phigma should roll out the feature?

My interpretation:

My hypothesis was not clearly supported. I predicted that the AI assistant would help newer users more, but the interaction coefficient was positive (`coef = 0.046`), which points in the opposite direction: the treatment effect may become slightly larger as prior whiteboard experience increases. However, the interaction p-value was about `0.098` , which is above the `0.05` significance threshold, so I would not make a strong claim that the treatment effect truly differs by prior use. The grouped bar chart still shows that treatment users generally have longer sessions across low, medium, and high prior-use groups. Based on this, the safest interpretation is that the AI assistant appears broadly useful, but this dataset does not give strong evidence that beginners or power users benefit more. For rollout, I would recommend making the feature available widely, while continuing to collect more data before using prior whiteboard experience as a targeting rule.

Part B — Logistic Regression: Who Converts?

Dataset: `phigma_subscription.csv` **Research question:** Which behavioral and demographic factors predict whether a free-tier user subscribes within 30 days?

This dataset requires **feature engineering before modeling**. Several columns arrive as raw datetimes or aggregate counts that need transformation.

```
In [16]: df_sub = pd.read_csv('phigma_subscription.csv')

print(f"Shape: {df_sub.shape}")
print(f"\nColumn types:\n{df_sub.dtypes}")
df_sub.head(3)
```

Shape: (3000, 15)

Column types:

user_id	object
signup_date	object
first_session_date	object
device_type	object
country_tier	object
referral_source	object
num_sessions_30d	int64
total_time_min_30d	int64
projects_created_30d	int64
exports_30d	int64
collaborators_invited_30d	int64
support_tickets_30d	int64
used_premium_feature_preview	int64
ai_assistant_sessions_30d	int64
subscribed	int64
dtype:	object

Out[16]:

	user_id	signup_date	first_session_date	device_type	country_tier	referral_source	num_sessions_30d	total_time_min_30d	projec
0	V94735	2024-05-23	2024-05-28	tablet	tier_1	organic	1	96	
1	V15501	2024-03-05	2024-03-05	tablet	tier_2	social	15	267	
2	V98542	2024-01-05	2024-01-19	desktop	tier_1	organic	16	263	

```
In [17]: df_sub.describe(include='all').round(2)
```

Out [17]:	user_id	signup_date	first_session_date	device_type	country_tier	referral_source	num_sessions_30d	total_time_min_30d
count	3000	3000	3000	3000	3000	3000	3000.00	3000.00
unique	2947	181	200	3	3	4	NaN	NaN
top	V43688	2024-06-06	2024-06-27	desktop	tier_1	organic	NaN	NaN
freq	3	33	29	1788	1332	1153	NaN	NaN
mean	NaN	NaN	NaN	NaN	NaN	NaN	9.25	216.17
std	NaN	NaN	NaN	NaN	NaN	NaN	5.09	139.20
min	NaN	NaN	NaN	NaN	NaN	NaN	1.00	5.00
25%	NaN	NaN	NaN	NaN	NaN	NaN	6.00	115.75
50%	NaN	NaN	NaN	NaN	NaN	NaN	8.00	188.50
75%	NaN	NaN	NaN	NaN	NaN	NaN	12.00	285.00
max	NaN	NaN	NaN	NaN	NaN	NaN	33.00	1004.00

B1 — Feature Engineering

Derive the required features listed in the README before doing any modeling or EDA.

```
In [18]: # Feature engineering for the subscription dataset.
# These transformations follow the lecture ideas: dates need to become numeric features,
# categories need encoding, and ratios need safe handling.

# 1. Parse date columns
_df = df_sub.copy()
_df['signup_date'] = pd.to_datetime(_df['signup_date'])
_df['first_session_date'] = pd.to_datetime(_df['first_session_date'])

# Save back to df_sub so later EDA can use the engineered columns.
df_sub = _df

# 2. Days between signup and first session
# A smaller value means the user became active quickly after signup.
df_sub['days_to_first_session'] = (
    df_sub['first_session_date'] - df_sub['signup_date']
).dt.days
```

```

# 3. Average session duration.
# I explicitly handle the divide-by-zero edge case even though this dataset has no zero-session rows.
df_sub['avg_session_duration_min'] = np.where(
    df_sub['num_sessions_30d'] == 0,
    0,
    df_sub['total_time_min_30d'] / df_sub['num_sessions_30d']
)

# 4. Weekend signup indicator
# Monday=0, ..., Saturday=5, Sunday=6
df_sub['is_weekend_signup'] = df_sub['signup_date'].dt.dayofweek.isin([5, 6]).astype(int)

# 5. Country tier is ordinal because tier_1, tier_2, and tier_3 have a natural ranking.
df_sub['country_tier_enc'] = df_sub['country_tier'].map({
    'tier_1': 1,
    'tier_2': 2,
    'tier_3': 3
})

# 6. One-hot encode unordered categories.
device_dummies_sub = pd.get_dummies(df_sub['device_type'], prefix='device', drop_first=True, dtype=int)
referral_dummies_sub = pd.get_dummies(df_sub['referral_source'], prefix='ref', drop_first=True, dtype=int)

# I use average session duration instead of total_time_min_30d in the model.
# total_time_min_30d is mechanically related to num_sessions_30d and avg_session_duration_min,
# so keeping all three would make the interpretation less clean.
raw_numeric = [
    'num_sessions_30d',
    'projects_created_30d',
    'exports_30d',
    'collaborators_invited_30d',
    'support_tickets_30d',
    'used_premium_feature_preview',
    'ai_assistant_sessions_30d'
]

engineered_features = [
    'days_to_first_session',
    'avg_session_duration_min',
    'is_weekend_signup',
    'country_tier_enc'
]

X_sub = pd.concat(
    [df_sub[raw_numeric + engineered_features], device_dummies_sub, referral_dummies_sub],

```

```

    axis=1
).astype(float)
y_sub = df_sub['subscribed'].astype(int)

# Quality checks
print(f"Zero-session users: {(df_sub['num_sessions_30d'] == 0).sum()}")
print(f"Any NaN values in X_sub: {X_sub.isna().any().any()}")
print(f"Any infinite values in X_sub: {np.isinf(X_sub.to_numpy()).any()}")
print(f"All features numeric: {all(X_sub.dtypes != 'object')}")
print(f"Feature matrix shape: {X_sub.shape}")
print('\nFeature names:')
print(list(X_sub.columns))

X_sub.head(3)

```

```

Zero-session users: 0
Any NaN values in X_sub: False
Any infinite values in X_sub: False
All features numeric: True
Feature matrix shape: (3000, 16)

```

```

Feature names:
['num_sessions_30d', 'projects_created_30d', 'exports_30d', 'collaborators_invited_30d', 'support_tickets_30d', 'used_premium_feature_preview', 'ai_assistant_sessions_30d', 'days_to_first_session', 'avg_session_duration_min', 'is_weekend_signup', 'country_tier_enc', 'device_mobile', 'device_tablet', 'ref_paid_ad', 'ref_referral', 'ref_social']

```

```

Out[18]:
  num_sessions_30d  projects_created_30d  exports_30d  collaborators_invited_30d  support_tickets_30d  used_premium_feature_pre
0                1.0                1.0          0.0                0.0                0.0
1                15.0                8.0          1.0                5.0                0.0
2                16.0                0.0          0.0                3.0                3.0

```

Your interpretation (B1):

For each engineered feature, one sentence: why might it predict subscription? Also describe how you handled the divide-by-zero edge case for `avg_session_duration_min`.

My interpretation:

`days_to_first_session` may predict subscription because users who start using the product soon after signup are likely more activated and more motivated. `avg_session_duration_min` may predict subscription because longer average sessions suggest deeper engagement, not just repeated short visits. `is_weekend_signup` may matter because weekend users might have different work habits,

personal-project goals, or urgency compared with weekday users. I encoded `country_tier` as an ordinal variable because the tiers have a natural market order, and moving from tier 1 to tier 3 may reflect differences in pricing, payment access, or market maturity. I one-hot encoded `device_type` and `referral_source` because they are unordered categories, so assigning them values like 1, 2, and 3 would create a fake ranking that the model could misread. I kept behavioral count features such as sessions, projects, exports, collaborators, support tickets, premium preview, and AI assistant sessions because they describe how actively the user used Phigma during the first 30 days. For the divide-by-zero edge case, I used `np.where()` so that users with zero sessions would receive an average session duration of `0` instead of producing an infinite or missing value; this dataset had no zero-session users, but the code still handles the case safely.

B2 — EDA for Subscription

```
In [19]: # Subscription rate
rate = df_sub['subscribed'].mean()
n_pos = df_sub['subscribed'].sum()
n_total = len(df_sub)
print(f"Subscription rate: {rate:.1%} ({n_pos} / {n_total})")
print(f"Class balance - 0: {n_total - n_pos} | 1: {n_pos}")
```

```
Subscription rate: 20.8% (625 / 3000)
Class balance - 0: 2375 | 1: 625
```

```
In [20]: # Visualize subscription relationships for more than three features.
# I use simple bar charts and boxplots because they match the EDA style from the lectures.

fig, axes = plt.subplots(2, 2, figsize=(13, 9))

# 1. Subscription rate by referral source
sub_by_ref = df_sub.groupby('referral_source')['subscribed'].mean().sort_values(ascending=False)
axes[0, 0].bar(sub_by_ref.index, sub_by_ref.values, color=ACCENT, edgecolor='black')
axes[0, 0].set_title('Subscription rate by referral source')
axes[0, 0].set_xlabel('Referral source')
axes[0, 0].set_ylabel('Subscription rate')
axes[0, 0].set_ylim(0, 0.45)
axes[0, 0].tick_params(axis='x', rotation=20)
for i, value in enumerate(sub_by_ref.values):
    axes[0, 0].text(i, value + 0.015, f'{value:.1%}', ha='center', fontsize=9)

# 2. Subscription rate by premium feature preview
sub_by_preview = df_sub.groupby('used_premium_feature_preview')['subscribed'].mean()
preview_labels = ['No preview', 'Previewed']
axes[0, 1].bar(preview_labels, sub_by_preview.values, color=ACCENT, edgecolor='black')
axes[0, 1].set_title('Subscription rate by premium preview')
axes[0, 1].set_xlabel('Premium feature preview')
```



```





axes[0, 1].set_ylabel('Subscription rate')
axes[0, 1].set_ylim(0, 0.45)
for i, value in enumerate(sub_by_preview.values):
    axes[0, 1].text(i, value + 0.015, f'{value:.1%}', ha='center', fontsize=9)

# 3. Subscription rate by country tier
sub_by_tier = df_sub.groupby('country_tier')['subscribed'].mean().sort_index()
axes[1, 0].bar(sub_by_tier.index, sub_by_tier.values, color=ACCENT, edgecolor='black')
axes[1, 0].set_title('Subscription rate by country tier')
axes[1, 0].set_xlabel('Country tier')
axes[1, 0].set_ylabel('Subscription rate')
axes[1, 0].set_ylim(0, 0.45)
for i, value in enumerate(sub_by_tier.values):
    axes[1, 0].text(i, value + 0.015, f'{value:.1%}', ha='center', fontsize=9)

# 4. Exports by subscription status
sns.boxplot(data=df_sub, x='subscribed', y='exports_30d', ax=axes[1, 1], color=ACCENT)
axes[1, 1].set_title('Exports by subscription status')
axes[1, 1].set_xlabel('Subscribed')
axes[1, 1].set_ylabel('Exports in first 30 days')
axes[1, 1].set_xticklabels(['No', 'Yes'])

plt.tight_layout()
plt.show()

```



```
In [21]: # Correlation matrix for numeric behavioral features.
# I include total_time_min_30d here for EDA so I can see its correlation with sessions,
# even though I did not include total_time_min_30d in the final model.
```

```
eda_numeric = [
    'num_sessions_30d',
```

```

    'total_time_min_30d',
    'avg_session_duration_min',
    'projects_created_30d',
    'exports_30d',
    'collaborators_invited_30d',
    'support_tickets_30d',
    'ai_assistant_sessions_30d',
    'days_to_first_session'
]


corr = df_sub[eda_numeric].corr()

fig, ax = plt.subplots(figsize=(10, 8))
sns.heatmap(corr, annot=True, fmt='.2f', cmap='coolwarm', linewidths=0.4, ax=ax)
ax.set_title('Correlation matrix for numeric features')
plt.tight_layout()
plt.show()

# Print the strongest correlations, excluding a variable with itself.
pairs = corr.abs().unstack().sort_values(ascending=False)
pairs = pairs[pairs < 1]
print('Strongest absolute correlations:')
print(pairs.head(8).round(3))

```

Correlation matrix for numeric features



Strongest absolute correlations:

total_time_min_30d	num_sessions_30d	0.843
num_sessions_30d	total_time_min_30d	0.843
avg_session_duration_min	total_time_min_30d	0.417
total_time_min_30d	avg_session_duration_min	0.417
ai_assistant_sessions_30d	projects_created_30d	0.036
projects_created_30d	ai_assistant_sessions_30d	0.036
exports_30d	ai_assistant_sessions_30d	0.027
ai_assistant_sessions_30d	exports_30d	0.027

dtype: float64

In [22]: *# This table gives a quick numeric summary of how selected features differ by subscription status.*

```
summary_by_sub = (
    df_sub
    .groupby('subscribed')[eda_numeric + ['used_premium_feature_preview']]
    .agg(['mean', 'median'])
    .round(2)
)
display(summary_by_sub)

# Simple correlations with the binary outcome are not a full model,
# but they are helpful for EDA.
corr_cols = eda_numeric + ['used_premium_feature_preview', 'country_tier_enc', 'subscribed']
outcome_corr = (
    df_sub[corr_cols]
    .corr()['subscribed']
    .drop('subscribed')
    .sort_values(key=abs, ascending=False)
)

print('Correlations with subscribed:')
print(outcome_corr.round(3))
```

	num_sessions_30d		total_time_min_30d		avg_session_duration_min		projects_created_30d		exports_30d		collaborators_invited
	mean	median	mean	median	mean	median	mean	median	mean	median	mean
subscribed											
0	9.15	8.0	211.83	182.0	23.18	23.00	2.43	2.0	1.84	1.0	1.80
1	9.63	9.0	232.67	202.0	24.61	24.75	2.43	2.0	2.66	2.0	2.09

```
Correlations with subscribed:
used_premium_feature_preview    0.276
exports_30d                      0.163
days_to_first_session          -0.115
country_tier_enc                 -0.102
avg_session_duration_min        0.072
total_time_min_30d              0.061
collaborators_invited_30d       0.056
support_tickets_30d             0.055
num_sessions_30d                0.038
ai_assistant_sessions_30d       0.020
projects_created_30d            -0.000
Name: subscribed, dtype: float64
```

Your interpretation (B2):

Which two or three features appear most strongly associated with subscription in your EDA? Are any findings counterintuitive? Note any features that appear highly correlated with each other — does this concern you?

My interpretation:

The overall subscription rate is about **20.8%**, so the dataset is imbalanced because most users did not subscribe. The strongest EDA association is **used_premium_feature_preview**: users who previewed a premium feature subscribed at a much higher rate than users who did not. **exports_30d** also appears important because subscribers tend to export more, which makes sense because exporting means the user is producing work that may need to be saved, shared, or used outside the app. **days_to_first_session** appears negatively associated with subscription, meaning users who wait longer before their first session are less likely to subscribe. One counterintuitive pattern is that **support_tickets_30d** is positively related to subscription; I interpret this as a possible sign that serious users contact support because they are trying to complete real work, not because they are simply unhappy. The main correlation concern is that **total_time_min_30d** and **num_sessions_30d** are strongly correlated ($r \approx 0.84$), so I did not include total time in the final model together with session count and average duration because it would make the model harder to interpret.

B3 — Logistic Regression with statsmodels (Inferential)

We fit on the **full dataset** here to maximize statistical power for inference. No train/test split — the goal is to interpret coefficients, not evaluate prediction.

```
In [23]: # Inferential logistic regression with statsmodels.
# This model uses the full dataset because the goal is coefficient interpretation.

X_const_sub = sm.add_constant(X_sub).astype(float)
```

```
model_logit = sm.Logit(y_sub, X_const_sub).fit(maxiter=200)
print(model_logit.summary())
```

Optimization terminated successfully.

Current function value: 0.440669

Iterations 6

Logit Regression Results

```
=====
Dep. Variable:          subscribed    No. Observations:          3000
Model:                  Logit        Df Residuals:              2983
Method:                 MLE          Df Model:                  16
Date:                  Mon, 06 Jul 2026    Pseudo R-squ.:          0.1389
Time:                  06:38:41          Log-Likelihood:         -1322.0
converged:              True            LL-Null:                 -1535.2
Covariance Type:        nonrobust        LLR p-value:             1.039e-80
=====
```

	coef	std err	z	P> z	[0.025	0.975]
const	-2.3413	0.259	-9.036	0.000	-2.849	-1.833
num_sessions_30d	0.0179	0.009	1.901	0.057	-0.001	0.036
projects_created_30d	-0.0096	0.021	-0.448	0.654	-0.052	0.032
exports_30d	0.1944	0.022	8.891	0.000	0.152	0.237
collaborators_invited_30d	0.0610	0.022	2.741	0.006	0.017	0.105
support_tickets_30d	0.1563	0.043	3.631	0.000	0.072	0.241
used_premium_feature_preview	1.4677	0.100	14.712	0.000	1.272	1.663
ai_assistant_sessions_30d	0.0202	0.017	1.205	0.228	-0.013	0.053
days_to_first_session	-0.0922	0.014	-6.714	0.000	-0.119	-0.065
avg_session_duration_min	0.0233	0.006	3.881	0.000	0.012	0.035
is_weekend_signup	0.0461	0.107	0.429	0.668	-0.165	0.257
country_tier_enc	-0.3557	0.065	-5.438	0.000	-0.484	-0.227
device_mobile	-0.1452	0.143	-1.013	0.311	-0.426	0.136
device_tablet	-0.2112	0.118	-1.794	0.073	-0.442	0.020
ref_paid_ad	-0.1817	0.136	-1.336	0.182	-0.448	0.085
ref_referral	0.2128	0.150	1.415	0.157	-0.082	0.507
ref_social	0.0560	0.122	0.458	0.647	-0.184	0.296

```
=====
```

In [24]: *# Convert logit coefficients into odds ratios.*
This follows the classification lecture: exp(beta) gives the multiplicative change in odds.

```
odds_ratios = np.exp(model_logit.params)
conf_int = np.exp(model_logit.conf_int())
conf_int.columns = ['OR_lower_95', 'OR_upper_95']

or_table = pd.DataFrame({
    'Odds Ratio': odds_ratios,
```

```

    '95% CI Lower': conf_int['OR_lower_95'],
    '95% CI Upper': conf_int['OR_upper_95'],
    'p-value': model_logit.pvalues
})

significant_or = (
    or_table
    .drop(index='const')
    .query('`p-value` < 0.05')
    .sort_values('Odds Ratio', ascending=False)
)

print('Significant predictors (p < 0.05):')
display(significant_or.round(4))

```

Significant predictors (p < 0.05):

	Odds Ratio	95% CI Lower	95% CI Upper	p-value
used_premium_feature_preview	4.3394	3.5687	5.2766	0.0000
exports_30d	1.2146	1.1636	1.2677	0.0000
support_tickets_30d	1.1692	1.0746	1.2722	0.0003
collaborators_invited_30d	1.0629	1.0175	1.1103	0.0061
avg_session_duration_min	1.0236	1.0116	1.0357	0.0001
days_to_first_session	0.9119	0.8877	0.9368	0.0000
country_tier_enc	0.7007	0.6164	0.7965	0.0000

In [25]: *# Visualize odds ratios. The vertical line at 1 means no change in odds.*

```

or_all = or_table.drop(index='const').sort_values('Odds Ratio')

fig, ax = plt.subplots(figsize=(8, 7))
colors = [ACCENT if p < 0.05 else '#cccccc' for p in or_all['p-value']]
ax.barh(or_all.index, or_all['Odds Ratio'], color=colors, edgecolor='black')

# Confidence interval lines
for i, (_, row) in enumerate(or_all.iterrows()):
    ax.plot([row['95% CI Lower'], row['95% CI Upper']], [i, i], color='black', linewidth=1)


ax.axvline(1, color='black', linestyle='--', linewidth=1)
ax.set_xscale('log')

```



```
ax.set_xlabel('Odds ratio (log scale)')  
ax.set_title('Logistic regression odds ratios')  
plt.tight_layout()  
plt.show()
```

Logistic regression odds ratios



 Results paragraph (B3):

Interpret at least three odds ratios in plain English (e.g. "Users who previewed a premium feature were X times more likely to subscribe..."). Identify at least one predictor whose sign or magnitude surprised you and offer an explanation.

My results paragraph:

The strongest predictor of subscription was `used_premium_feature_preview`. Holding the other variables constant, users who previewed a premium feature had about `4.34` times the odds of subscribing compared with users who did not preview one. `exports_30d` was also significant: each additional export was associated with about `1.21` times the odds of subscription, which suggests that users who produce downloadable or shareable work are more likely to see value in paying. `days_to_first_session` had an odds ratio of about `0.91`, meaning each extra day of delay before the first session reduced the odds of subscription by about 9%. `country_tier_enc` was also significant with an odds ratio below 1, suggesting that users in higher-numbered country tiers had lower odds of subscribing after controlling for behavior. A nuanced finding is that `support_tickets_30d` was positive and significant; at first this seems surprising because support problems sound negative, but it may mean that serious users submit tickets because they are actively trying to use the product for real work. `ai_assistant_sessions_30d` was not significant after controlling for other variables, so in this model AI usage alone appears less predictive than clearer purchase-intent signals such as premium preview and exports.

B4 — Logistic Regression with sklearn (Predictive)

In [26]: *# Train/test split for prediction. Stratify keeps the positive subscription rate similar
in both training and test data, which matters because the classes are imbalanced.*

```
X_train, X_test, y_train, y_test = train_test_split(
    X_sub,
    y_sub,
    test_size=0.2,
    random_state=42,
    stratify=y_sub
)

print(f"Train shape: {X_train.shape}")
print(f"Test shape: {X_test.shape}")
print(f"Train positive rate: {y_train.mean():.3f}")
print(f"Test positive rate: {y_test.mean():.3f}")
```

Train shape: (2400, 16)
Test shape: (600, 16)
Train positive rate: 0.208
Test positive rate: 0.208

```
In [27]: # Fit sklearn LogisticRegression.  
# I standardize the features because logistic regression is easier to optimize when features are on similar scales.  
  
scaler = StandardScaler()  
X_train_scaled = scaler.fit_transform(X_train)  
X_test_scaled = scaler.transform(X_test)  
  
clf_lr = LogisticRegression(max_iter=1000, random_state=42)  
clf_lr.fit(X_train_scaled, y_train)  
  
y_pred_lr = clf_lr.predict(X_test_scaled)  
y_prob_lr = clf_lr.predict_proba(X_test_scaled)[:, 1]  
  
print('Logistic regression model fitted.')
```

Logistic regression model fitted.

```
In [28]: # Confusion matrix for logistic regression.  
  
cm_lr = confusion_matrix(y_test, y_pred_lr)  
disp = ConfusionMatrixDisplay(  
    confusion_matrix=cm_lr,  
    display_labels=['Not subscribed', 'Subscribed']  
)  
  
fig, ax = plt.subplots(figsize=(5, 4))  
disp.plot(ax=ax, cmap='Blues', colorbar=False)  
ax.set_title('Logistic Regression - Confusion Matrix')  
plt.tight_layout()  
plt.show()  
  
tn_lr, fp_lr, fn_lr, tp_lr = cm_lr.ravel()  
print(f"TN: {tn_lr} | FP: {fp_lr} | FN: {fn_lr} | TP: {tp_lr}")
```

Logistic Regression — Confusion Matrix



TN: 459 | FP: 16 | FN: 103 | TP: 22

In [29]: *# Classification metrics and ROC-AUC.*

```
print(classification_report(y_test, y_pred_lr, target_names=['Not subscribed', 'Subscribed'], zero_division=0))
auc_lr = roc_auc_score(y_test, y_prob_lr)
print(f"ROC-AUC: {auc_lr:.4f}")
```

Save report dictionary for Part C comparison.

```
lr_report = classification_report(y_test, y_pred_lr, output_dict=True, zero_division=0)
```


	precision	recall	f1-score	support
Not subscribed	0.82	0.97	0.89	475
Subscribed	0.58	0.18	0.27	125
accuracy			0.80	600
macro avg	0.70	0.57	0.58	600
weighted avg	0.77	0.80	0.76	600

ROC-AUC: 0.7373

In [30]: *# ROC curve for logistic regression.*

```
fpr_lr, tpr_lr, _ = roc_curve(y_test, y_prob_lr)

fig, ax = plt.subplots(figsize=(6, 5))
ax.plot(fpr_lr, tpr_lr, color=ACCENT, linewidth=2, label=f'Logistic Regression (AUC = {auc_lr:.3f})')
ax.plot([0, 1], [0, 1], color='black', linestyle='--', linewidth=1, label='Random baseline')
ax.set_xlabel('False Positive Rate')
ax.set_ylabel('True Positive Rate')
ax.set_title('ROC Curve - Logistic Regression')
ax.legend()
plt.tight_layout()
plt.show()
```



 **Your interpretation (B4):**

Which metric did you decide is most important given the class imbalance, and why? Interpret the confusion matrix: how many false negatives does the model produce? From Phigma's business perspective, what is the cost of a false negative (a user who would have subscribed but the model missed)?

My interpretation:

Because only about 21% of users subscribed, accuracy is not the most informative metric. This connects directly to the classification lecture, where accuracy was shown to be misleading when one class is much more common than the other. In this dataset, a model could get a high accuracy score by mostly predicting the majority class, which is `not subscribed`, but that would not mean it is good at finding the users Phigma cares about. For that reason, I focus more on class-1 recall, class-1 F1, and ROC-AUC. Recall tells me how many actual subscribers the model catches, F1 balances recall and precision for the subscriber class, and ROC-AUC tells me how well the model ranks likely subscribers above unlikely subscribers across thresholds. The logistic regression model had an ROC-AUC of about 0.737, which means it has a moderate ability to separate likely subscribers from unlikely subscribers. The confusion matrix shows 459 true negatives, 16 false positives, 103 false negatives, and 22 true positives. This means the model missed 103 users who actually subscribed, which is a major limitation at the default 0.5 threshold. From Phigma's business perspective, false negatives matter because these are users with real conversion potential who might not receive targeted onboarding, sales messaging, or premium-feature nudges. If Phigma cares more about finding potential subscribers than avoiding extra outreach, I would consider lowering the decision threshold to improve recall, just like how we usually discussed when explaining the precision-recall tradeoff.

Part C — Your Model: Choose, Learn, Apply, Compare

Dataset: `phigma_subscription.csv` **Same train/test split as B4.** Do not re-split.

C1 — Choose Your Model

Available options: Decision Tree, Random Forest, Gradient Boosting, XGBoost, LightGBM, K-Nearest Neighbors, Naive Bayes, SVM, LDA, AdaBoost.

 **Your chosen model:** `Gaussian Naive Bayes`

Conceptual explanation (write this before any code):

3–5 sentences. How does this model learn from training data? What does it optimize? How is it structurally different from logistic regression? Why might it perform differently on this particular dataset given what you know about the features?

My explanation:

I chose Gaussian Naive Bayes because it connects directly to the probability, normal distribution, and Bayes' theorem ideas from the lectures. Bayes' theorem was used to update a prior belief after seeing evidence. Gaussian Naive Bayes applies that same idea to classification: it starts with the base rate of each class, then updates that probability based on the feature values for a user. The model learns

from the training data by estimating the mean and variance of each numeric feature separately within each class. Under the Gaussian version, it assumes each feature follows a bell-shaped distribution within each class, which connects to the normal distribution. It then chooses the class with the highest posterior probability. It is called **naive** because it assumes features are conditionally independent within each class, which is a strong simplification because user behavior features can be related to each other. This is structurally different from logistic regression because logistic regression learns one linear log-odds equation, while Naive Bayes builds separate class-conditional probability estimates. It might perform differently here because the dataset has clear probability signals, such as premium preview and exports, but the independence assumption may be imperfect because engagement features can still move together.

C2 — Fit and Evaluate

```
In [31]: # Fit Gaussian Naive Bayes using the same train/test split from B4.  
# I use the scaled matrices already created for logistic regression.
```


```
clf_nb = GaussianNB()  
clf_nb.fit(X_train_scaled, y_train)  
  
y_pred_nb = clf_nb.predict(X_test_scaled)  
y_prob_nb = clf_nb.predict_proba(X_test_scaled)[:, 1]  
  
print('Gaussian Naive Bayes model fitted.')
```

Gaussian Naive Bayes model fitted.

```
In [32]: # Confusion matrix for Gaussian Naive Bayes.
```

```
cm_nb = confusion_matrix(y_test, y_pred_nb)  
disp = ConfusionMatrixDisplay(  
    confusion_matrix=cm_nb,  
    display_labels=['Not subscribed', 'Subscribed']  
)  
  
fig, ax = plt.subplots(figsize=(5, 4))  
disp.plot(ax=ax, cmap='Blues', colorbar=False)  
ax.set_title('Gaussian Naive Bayes - Confusion Matrix')  
plt.tight_layout()  
plt.show()  
  
tn_nb, fp_nb, fn_nb, tp_nb = cm_nb.ravel()  
print(f"TN: {tn_nb} | FP: {fp_nb} | FN: {fn_nb} | TP: {tp_nb}")
```

Gaussian Naive Bayes — Confusion Matrix



TN: 436 | FP: 39 | FN: 79 | TP: 46

In [33]: *# Classification report and ROC-AUC for Gaussian Naive Bayes.*

```
print(classification_report(y_test, y_pred_nb, target_names=['Not subscribed', 'Subscribed'], zero_division=0))
auc_nb = roc_auc_score(y_test, y_prob_nb)
print(f"ROC-AUC: {auc_nb:.4f}")

nb_report = classification_report(y_test, y_pred_nb, output_dict=True, zero_division=0)
```


	precision	recall	f1-score	support
Not subscribed	0.85	0.92	0.88	475
Subscribed	0.54	0.37	0.44	125
accuracy			0.80	600
macro avg	0.69	0.64	0.66	600
weighted avg	0.78	0.80	0.79	600

ROC-AUC: 0.7478

In [34]: *# ROC curve for Gaussian Naive Bayes.*

```
fpr_nb, tpr_nb, _ = roc_curve(y_test, y_prob_nb)

fig, ax = plt.subplots(figsize=(6, 5))
ax.plot(fpr_nb, tpr_nb, color=ACCENT, linewidth=2, label=f'Gaussian NB (AUC = {auc_nb:.3f})')
ax.plot([0, 1], [0, 1], color='black', linestyle='--', linewidth=1, label='Random baseline')
ax.set_xlabel('False Positive Rate')
ax.set_ylabel('True Positive Rate')
ax.set_title('ROC Curve - Gaussian Naive Bayes')
ax.legend()
plt.tight_layout()
plt.show()
```



C3 — Model Comparison

In [35]: *# Summary comparison table for the positive class (subscribed = 1).*

```
comparison = pd.DataFrame({  
    'Model': ['Logistic Regression', 'Gaussian Naive Bayes'],  
    'ROC-AUC': [auc_lr, auc_nb],  
    'Precision (class 1)': [lr_report['1']['precision'], nb_report['1']['precision']],  
    'Recall (class 1)': [lr_report['1']['recall'], nb_report['1']['recall']],  
})
```

```

    'F1 (class 1)': [lr_report['1']['f1-score'], nb_report['1']['f1-score']],
    'False negatives': [fn_lr, fn_nb],
    'False positives': [fp_lr, fp_nb]
}).set_index('Model')

comparison.round(4)

```

Out [35]:

	ROC-AUC	Precision (class 1)	Recall (class 1)	F1 (class 1)	False negatives	False positives
Model						
Logistic Regression	0.7373	0.5789	0.176	0.2699	103	16
Gaussian Naive Bayes	0.7478	0.5412	0.368	0.4381	79	39


```

In [36]: # Plot both ROC curves on the same axes.

fig, ax = plt.subplots(figsize=(7, 5))
ax.plot(fpr_lr, tpr_lr, color=ACCENT, linewidth=2, label=f'Logistic Regression (AUC = {auc_lr:.3f})')
ax.plot(fpr_nb, tpr_nb, color='steelblue', linewidth=2, label=f'Gaussian NB (AUC = {auc_nb:.3f})')
ax.plot([0, 1], [0, 1], color='black', linestyle='--', linewidth=1, label='Random baseline')
ax.set_xlabel('False Positive Rate')
ax.set_ylabel('True Positive Rate')
ax.set_title('ROC Curve Comparison')
ax.legend()
plt.tight_layout()
plt.show()

```

ROC Curve Comparison



Comparison paragraph (C3):

*Which model performs better, and by how much? Explain *why* the performance difference occurred — based on what you know about how each model works and the structure of this dataset (non-linearities, class imbalance, number of features). If your model performed worse than logistic regression, that is a valid finding — explain why.**

My comparison paragraph:

Gaussian Naive Bayes performed slightly better than logistic regression on ROC-AUC and clearly better on class-1 recall and class-1 F1. Logistic regression had an ROC-AUC of about `0.737`, class-1 recall of about `0.176`, and class-1 F1 of about `0.270`. Gaussian Naive Bayes had an ROC-AUC of about `0.748`, class-1 recall of about `0.368`, and class-1 F1 of about `0.438`. The practical difference is that Naive Bayes caught more actual subscribers: logistic regression missed `103` subscribers, while Naive Bayes missed `79`. However, Naive Bayes also produced more false positives (`39` compared with `16`), so the improvement comes with a tradeoff. This tradeoff connects to the classification because changing how willing a model is to predict the positive class usually changes false positives and false negatives at the same time. I think this happened because the default logistic regression classifier was conservative in an imbalanced dataset. Logistic regression learns a linear log-odds boundary and then uses the default `0.5` threshold, so it only predicts subscription when the estimated probability is high enough. Gaussian Naive Bayes works differently because it combines class-conditional probability evidence from each feature, so several moderate signals can push a user toward the subscriber class even if the final decision creates more false positives. If Phigma wants to avoid bothering users who are unlikely to subscribe, logistic regression is safer; if Phigma wants to find more possible subscribers for outreach, Gaussian Naive Bayes is more useful in this test.

C4 — Bonus: Hyperparameter Tuning (+5 points)

Complete this section only if you want to attempt the bonus.

```
In [37]: # Bonus tuning for Gaussian Naive Bayes.
# GaussianNB has one main smoothing hyperparameter, so I test several values.

param_grid = {
    'var_smoothing': np.logspace(-12, -6, 7)
}

search = GridSearchCV(
    GaussianNB(),
    param_grid,
    cv=5,
    scoring='roc_auc',
    n_jobs=1
)
search.fit(X_train_scaled, y_train)

print(f"Best parameters: {search.best_params}")
print(f"Best CV AUC:      {search.best_score_:.4f}")

y_pred_tuned = search.best_estimator_.predict(X_test_scaled)
y_prob_tuned = search.best_estimator_.predict_proba(X_test_scaled)[:, 1]
auc_tuned = roc_auc_score(y_test, y_prob_tuned)
```

```

print(f"Tuned test AUC: {auc_tuned:.4f}")
print(f"Default test AUC: {auc_nb:.4f}")
print(classification_report(y_test, y_pred_tuned, target_names=['Not subscribed', 'Subscribed'], zero_division=0))

if abs(auc_tuned - auc_nb) < 0.005:
    print('Tuning did not meaningfully change performance; the default GaussianNB model was already close to the tuned
else:
    print('Tuning changed performance enough to be noticeable on the test set.')

```

Best parameters: {'var_smoothing': np.float64(1e-12)}

Best CV AUC: 0.7150

Tuned test AUC: 0.7478

Default test AUC: 0.7478

	precision	recall	f1-score	support
Not subscribed	0.85	0.92	0.88	475
Subscribed	0.54	0.37	0.44	125
accuracy			0.80	600
macro avg	0.69	0.64	0.66	600
weighted avg	0.78	0.80	0.79	600

Tuning did not meaningfully change performance; the default GaussianNB model was already close to the tuned result.

Submission Checklist

Before submitting, confirm:

- ☐ All `# YOUR CODE HERE` cells are completed and run without errors
- ☐ All `** Your interpretation**` cells contain your own written analysis
- ☐ The notebook runs top-to-bottom without errors (`Kernel > Restart & Run All`)
- ☐ File is named `iat461_a2_[YourName]_[SFUid].ipynb`
- ☐ PDF export is named `iat461_a2_[YourName]_[SFUid].pdf`
- ☐ All AI-assisted code blocks include `#BEGIN / #END` disclosure comments